



Algorithm Design VIII

Greedy Algorithms


Guoqiang Li
School of Software



SHANGHAI JIAO TONG
UNIVERSITY

Minimum Spanning Trees

Build a Network




Suppose you are asked to **network** a collection of computers by linking selected pairs of them.

This translates into a graph problem in which

- nodes are computers,
- undirected edges are potential links, each with a **maintenance cost**.

Build a Network



The goal is to

- pick enough of these edges that the nodes are **connected**,
- the total maintenance cost is **minimum**.

One immediate observation is that the optimal set of edges cannot contain a **cycle**.



Properties of the Optimal Solutions

Lemma (1)

Removing a cycle edge cannot **disconnect** a graph.

So the solution must be **connected** and **acyclic**: undirected graphs of this kind are called **trees**.

A tree with **minimum total weight**, is a **minimum spanning tree**, MST.

Input: An undirected graph $G = (V, E)$; edge weights w_e

Output: A tree $T = (V, E')$ with $E' \subseteq E$ that **minimizes**

$$\text{weight}(T) = \sum_{e \in E'} w_e$$



Lemma (2)

A tree on n nodes has $n - 1$ edges.

To build the tree one edge at a time, starting from an empty graph.

Each of the n nodes is disconnected from the others, in a connected component by itself.

As edges are added, these components merge. Since each edge unites two different components, exactly $n - 1$ edges are added by the time the tree is fully formed.

When a particular edge (u, v) comes up, we can be sure that u and v lie in separate connected components, for otherwise there would already be a path between them and this edge would create a cycle.



Lemma (3)

Any connected, undirected graph $G = (V, E)$ with $|E| = |V| - 1$ is a tree.

It is the converse of Lemma (2). We just need to show that G is acyclic.

While the graph contains a cycle, remove one edge from this cycle.

The process terminates with some graph $G' = (V, E')$, $E' \subseteq E$, which is acyclic and, by Lemma (1), is also connected.

Therefore G' is a tree, whereupon $|E'| = |V| - 1$ by Lemma (2). So $E' = E$, no edges were removed, and G was acyclic to start with.



Lemma (4)

An undirected graph is a **tree** if and only if there is a **unique** path between any pair of nodes.

In a tree, any two nodes can only have **one path** between them; for if there were two paths, the union of these paths would contain a cycle.

On the other hand, if a graph has a path between any two nodes, then it is **connected**. If these paths are **unique**, then the graph is also acyclic.

A Greedy Approach





Kruskal's minimum spanning tree algorithm starts with the **empty graph** and then selects edges from E according to the following rule.

Repeatedly add the next lightest edge that doesn't produce a cycle.

Example

Starting with an empty graph and then attempt to add edges in increasing order of weight

$B - C; C - D; B - D; C - F; D - F; E - F; A - D; A - B; C - E; A - C$



The Cut Property



Lemma

Suppose edges X are part of a MST of $G = (V, E)$. Pick any subset of nodes S for which X does not cross between S and $V \setminus S$, and let e be the *lightest edge* across this partition. Then

$$X \cup \{e\}$$


is part of some *MST*.



The Cut Property

A **cut** is any **partition** of the vertices into two groups, S and $V \setminus S$.

It is **safe** to add the **lightest edge** across any **cut**, provided X has no edges across the cut.



Proof of the Cut Property



Proof:

Edges X are part of some MST T ; if the new edge e also happens to be part of T , then there is nothing to prove.

So assume e is not in T . We will construct a different MST T' containing $X \cup \{e\}$ by altering T slightly, changing just one of its edges.

Add edge e to T . Since T is **connected**, it already has a path between the endpoints of e , so adding e creates a **cycle**.

This cycle must also have some other edge e' across the cut $(S, V \setminus S)$. If we now remove e'

$$T' = T \cup \{e\} \setminus \{e'\}$$

which we will show to be a **tree**.

T' is connected by **Lemma (1)**, since e' is a cycle edge. And it has the same number of edges as T ; so by **Lemma (2)** and **Lemma (3)**, it is also a tree.

Proof of the Cut Property



Proof:

T' is a minimum spanning tree, since

$$\text{weight}(T') = \text{weight}(T) + w(e) - w(e')$$

Both e and e' cross between S and $V \setminus S$, and e is the lightest edge of this type. Therefore $w(e) \leq w(e')$, and


$$\text{weight}(T') \leq \text{weight}(T)$$

Since T is an MST, it must be the case that $\text{weight}(T') = \text{weight}(T)$ and that T' is also an MST.



An Example of Cut Property





(a)



(b)



(c)



Kruskal's Algorithm



KRUSKAL (G, w)

input : A connected undirected graph $G = (V, E)$, with edge weight w_e

output: A minimum spanning tree defined by the edges X

for *all* $u \in V$ **do**

 makeset (u);

end

$X = \{ \}$;

Sort the edges E by weight;

for *all* $(u, v) \in E$ *in increasing order of weight* **do**

if find (u) \neq find (v) **then**

 add (u, v) to X ;

 union (u, v)

end

end

Data Structure Retailer: Disjoint Sets



SHANGHAI JIAO TONG
UNIVERSITY

<code>makeset(x)</code>	create a singleton set containing x	$ V $
<code>find(x)</code>	find the set that x belong to	$2 \cdot E $
<code>union(x, y)</code>	merge the sets containing x and y	$ V - 1$

Prim's Algorithm

A General Kruskal's Algorithm



```
 $X = \{ \};$   
repeat until  $|X| = |V| - 1;$   
    pick a set  $S \subset V$  for which  $X$  has no edges between  $S$  and  
     $V - S;$   
    let  $e \in E$  be the minimum-weight edge between  $S$  and  $V - S;$   
     $X = X \cup \{e\};$ 
```

Prim's Algorithm



A popular alternative to **Kruskal's** algorithm is **Prim's**, in which the intermediate set of edges X always forms a subtree, and S is chosen to be the set of this tree's vertices.

On each iteration, the subtree defined by X grows by one edge.

The lightest edge between a vertex in S and a vertex outside S . We can equivalently think of S as growing to include the vertex $v \notin S$ of smallest **cost**:

$$\text{cost}(v) = \min_{u \in S} w(u, v)$$

The Algorithm



PRIM(G, w)

input : A connected undirected graph $G = (V, E)$, with edge weights w_e

output: A minimum spanning tree defined by the array $prev$

for *all* $u \in V$ **do**

$cost(u) = \infty$;

$prev(u) = nil$;

end

pick any initial node u_0 ;

$cost(u_0) = 0$;

$H = \text{makequeue}(V) \setminus \setminus$ *using cost-values as keys*;

while H *is not empty* **do**

$v = \text{deletemin}(H)$;

for *each* $(v, z) \in E$ **do**

if $cost(z) > w(v, z)$ **then**

$cost(z) = w(v, z)$; $prev(z) = v$;

$\text{decreasekey}(H, z)$;

end

end

end

Dijkstra's Algorithm



DIJKSTRA (G, l, s)

input : Graph $G = (V, E)$, directed or undirected; positive edge length $\{l_e \mid e \in E\}$;
Vertex $s \in V$

output: For all vertices u reachable from s , $dist(u)$ is the set to the distance from s to u

for all $u \in V$ **do**

$dist(u) = \infty$;
 $prev(u) = nil$;

end

$dist(s) = 0$;

$H = \text{makequeue}(V) \setminus \setminus$ *using dist-values as keys*;

while H is not empty **do**

$u = \text{deletemin}(H)$;

for all edge $(u, v) \in E$ **do**

if $dist(v) > dist(u) + l(u, v)$ **then**

$dist(v) = dist(u) + l(u, v)$; $prev(v) = u$;
 decreasekey (H, v);

end

end

end

Think About



Let C be a cycle with no red edges, and select an uncolored edge of C of max cost and color it red.

Let D be a edge set crossing a cut with no blue edges, and select an uncolored edge in D of min cost and color it blue.

Apply the red and blue rules nondeterministically until all edges are colored.

The blue edges form an MST.

Set Cover

The Problem




A county is in its early stages of planning and is deciding where to put schools.

There are only two constraints:


- each school should be **in a town**,
- and no one should have to travel more than **30** miles to reach one of them.

Q: What is the minimum number of schools needed?

(a)



(b)



The Problem



This is a typical (cardinality) set cover problem.

- For each town x , let S_x be the set of towns within 30 miles of it.
- A school at x will essentially “cover” these other towns.
- The question is then, how many sets S_x must be picked in order to cover all the towns in the county?


Set Cover Problem



SET COVER

- **Input:** A set of elements B , sets $S_1, \dots, S_m \subseteq B$
- **Output:** A selection of the S_i whose union is B .
- **Cost:** Number of sets picked.

The Example





Lemma

Suppose B contains n elements and that the *optimal cover* consists of OPT sets. Then the *greedy algorithm* will use at most $\ln n \cdot OPT$ sets.

Proof.

Let n_t be the number of elements still not covered after t iterations of the greedy algorithm (so $n_0 = n$).

Since these remaining elements are covered by the optimal OPT sets, there must be some set with at least n_t/OPT of them.

Therefore, the greedy strategy will ensure that

$$n_{t+1} \leq n_t - \frac{n_t}{OPT} = n_t \left(1 - \frac{1}{OPT}\right)$$

which by repeated application implies

$$n_t \leq n_0 \left(1 - \frac{1}{OPT}\right)^t$$



A more convenient bound can be obtained from the useful inequality

$$1 - x \leq e^{-x} \text{ for all } x$$

with equality if and only if $x = 0$,

Thus

$$n_t \leq n_0 \left(1 - \frac{1}{OPT}\right)^t < n_0 \left(e^{-\frac{1}{OPT}}\right)^t = n e^{-\frac{t}{OPT}}$$

At $t = \ln n \cdot OPT$, therefore, n_t is strictly less than $n e^{-\ln n} = 1$, which means no elements remain to be covered.

Why Greedy Does Not Work: Coin Changing

Coin changing



Goal. Given U. S. currency denominations $\{1, 5, 10, 25, 100\}$, devise a method to pay amount to customer using fewest coins.

Example \$34.

Cashier's algorithm. At each iteration, add coin of the largest value that does not take us past the amount to be paid.

Example \$2.89.

Cashier's algorithm



```
CASHIERS-ALGORITHM( $x, c_1, c_2, \dots, c_n$ )  
SORT  $n$  coin denominations so that  $0 < c_1 < c_2 < \dots < c_n$  ;  
 $S \leftarrow \emptyset$ ;  
while  $x > 0$  do  
     $k \leftarrow$  largest coin denomination  $c_k$  such that  $c_k \leq x$ ;  
    if no such  $k$  then RETURN no solution;  
    else  
         $x \leftarrow x - c_k$ ;  
         $S \leftarrow S \cup \{k\}$ ;  
    end  
end  
RETURN  $S$ ;
```




Is the cashier's algorithm optimal?

- A** Yes, greedy algorithms are always optimal.
- B** Yes, for any set of coin denominations $c_1 < c_2 < \dots < c_n$ provided $c_1 = 1$.
- C** Yes, because of special properties of U.S. coin denominations.
- D** No.

Cashier's algorithm (for arbitrary coin denominations)



Q. Is cashier's algorithm optimal for any set of denominations?

A. No. Consider U.S. postage: 1, 10, 21, 34, 70, 100, 350, 1225, 1500.

- Cashier's algorithm: $\$140 = 100 + 34 + 1 + 1 + 1 + 1 + 1 + 1$.
- Optimal: $\$140 = 70 + 70$.

A. No. It may not even lead to a feasible solution if $c_1 > 1$: 7, 8, 9.

- Cashier's algorithm: $\$15 = 9 + ?$.
- Optimal: $\$15 = 7 + 8$.

Properties of any optimal solution (for U.S. coin denominations)



SHANGHAI JIAO TONG
UNIVERSITY

Property. Number of pennies ≤ 4 .

Proof. Replace 5 pennies with 1 nickel.

Property. Number of nickels ≤ 1 .

Property. Number of quarters ≤ 3 .

Property. Number of nickels + number of dimes ≤ 2 .

Proof.

- Recall: ≤ 1 nickel.
- Replace 3 dimes and 0 nickels with 1 quarter and 1 nickel;
- Replace 2 dimes and 1 nickel with 1 quarter.

Optimality of cashier's algorithm (for U.S. coin denominations)



SHANGHAI JIAO TONG
UNIVERSITY

Theorem

Cashier's algorithm is optimal for U.S. coins $\{1, 5, 10, 25, 100\}$.



A rather formal proof

Proof. by induction on amount to be paid x

Consider optimal way to change $c_k \leq x \leq c_{k+1}$: greedy takes coin k .

Claim that any optimal solution must take coin k .

- if not, it needs enough coins of type c_1, \dots, c_{k-1} to add up to x .
- table below indicates no optimal solution can do this

Problem reduces to coin-changing $x - c_k$ cents, which, **by induction**, is optimally solved by cashier's algorithm.

k	c_k	all optimal solutions must satisfy	max value of c_1, c_2, \dots, c_{k-1} in any optimal solution
1	1	$P \leq 4$	none
2	5	$N \leq 1$	4
3	10	$N + D \leq 2$	$4 + 5 = 9$
4	25	$Q \leq 3$	$20 + 4 = 24$
5	100	no limit	$75 + 24 = 99$